

Task Documentation

This Document shows step-by-step task done with screenshot

Infrastructure Provisioning (Terraform)

Prerequisites

I already have the following:

- AWS Account
- AWS CLI installed in my local machine. Run to check: `aws --version`
- Terraform installed in my local machine. Run to check: `terraform version`

```
KRIS@Christian MINGW64 ~/Desktop/t4dev-capstone
● $ aws --version
aws-cli/2.8.12 Python/3.9.11 Windows/10 exe/AMD64 prompt/off

KRIS@Christian MINGW64 ~/Desktop/t4dev-capstone
● $ terraform version
Terraform v1.14.5
on windows_amd64

KRIS@Christian MINGW64 ~/Desktop/t4dev-capstone
○ $ [ ]
```

In my AWS Account, I created an access key and secret key for terraform to use.

The screenshot shows the 'Create access key' wizard in the AWS IAM console. The current step is 'Retrieve access keys'. It displays two fields: 'Access key' (containing 'AKIAUJPMYNOFLGPMWSXI') and 'Secret access key' (containing a long string of characters). A note above the fields states: 'This is the only time that the secret access key can be viewed or downloaded. You cannot recover it later. However, you can create a new access key any time.' On the left, a sidebar lists steps: Step 1 (selected), Step 2 - optional, Step 3, and 'Retrieve access keys' (which is also selected). At the bottom right are 'Download .csv file' and 'Done' buttons.

Then I configure the aws cli to use the access key and secret key.

```
aws configure
```

```
KRIS@Christian MINGW64 ~/Desktop/t4dev-capstone
● $ aws configure
AWS Access Key ID [*****WSXI]: ]]
AWS Secret Access Key [*****mhr/]:
Default region name [eu-west-2]:
Default output format [None]: ]]

KRIS@Christian MINGW64 ~/Desktop/t4dev-capstone
● $ aws sts get-caller-identity
{
    "UserId": "AIDAUPMYNOFL5UGJLLUH",
    "Account": "307946680662",
    "Arn": "arn:aws:iam::307946680662:user/tech4dev_group4"
} ]]

KRIS@Christian MINGW64 ~/Desktop/t4dev-capstone
○ $ ]]
```

Steps

1. Goto the terraform folder and Initialize Terraform

```
cd infrastructure/terraform
terraform init
```

```
KRIS@Christian MINGW64 ~/Desktop/t4dev-capstone
● $ cd infrastructure/terraform

KRIS@Christian MINGW64 ~/Desktop/t4dev-capstone/infrastructure/terraform
$

KRIS@Christian MINGW64 ~/Desktop/t4dev-capstone/infrastructure/terraform
● $ terraform init
Initializing the backend...
Initializing provider plugins...
- Finding hashicorp/aws versions matching "~> 6.0"...
- Installing hashicorp/aws v6.32.1...
- Installed hashicorp/aws v6.32.1 (signed by HashiCorp)
Terraform has created a lock file .terraform.lock.hcl to record the provider
selections it made above. Include this file in your version control repository
so that Terraform can guarantee to make the same selections by default when
you run "terraform init" in the future.

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.

KRIS@Christian MINGW64 ~/Desktop/t4dev-capstone/infrastructure/terraform
○ $ []
```

2. Apply Terraform `terraform apply`

```
aws_instance.worker: Still destroying... [id=i-069c0ec95b27986c8, 00m50s elapsed]
aws_instance.worker: Destruction complete after 58s

Apply complete! Resources: 2 added, 0 changed, 1 destroyed.

Outputs:

master_public_ip = "13.40.72.112"
vpc_id = "vpc-02558ba1002c1baf3"
worker-backend_public_ip = "35.178.102.235"
worker-frontend_public_ip = "3.8.216.19"

KRIS@Christian MINGW64 ~/Desktop/t4dev-capstone/infrastructure/terraform
○ $ []
```

The screenshot shows the AWS EC2 Instances page. It lists three instances:

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 IP	Elastic IP
k8s-worker-Backend	i-0f94037d22b601b20	Running	t3.small	3/3 checks passed	View alarms	eu-west-2a	ec2-35-178-102-235.eu...	35.178.102.235	-
k8s-master	i-063622f7c2d679aca	Running	t3.small	3/3 checks passed	View alarms	eu-west-2a	ec2-13-40-72-112.eu-w...	13.40.72.112	-
k8s-worker-frontend	i-085a5be71ceea8956	Running	t3.small	3/3 checks passed	View alarms	eu-west-2a	ec2-3-8-216-19.eu-wes...	3.8.216.19	-

Configuration setup (Ansible)

Prerequisites

I already have the following:

- Ansible installed in my local machine via wsl. Run to check: `ansible --version`

```
kris@Christian:~$ ansible --version
ansible [core 2.16.3]
  config file = None
  configured module search path = ['/home/kris/.ansible/plugins/modules', '/usr/share/ansible/plugins/modules']
  ansible python module location = /usr/lib/python3/dist-packages/ansible
  ansible collection location = /home/kris/.ansible/collections:/usr/share/ansible/collections
  executable location = /usr/bin/ansible
  python version = 3.12.3 (main, Nov  6 2025, 13:44:16) [GCC 13.3.0] (/usr/bin/python3)
  jinja version = 3.1.2
  libyaml = True
kris@Christian:~$
```

Steps

1. Goto the ansible folder and update the inventory file `cd infrastructure/ansible`

```
inventory.ini
[masters]
# Inputted my Master Node Public IP
k8s-master ansible_host=13.40.72.112 ansible_user=ubuntu ansible_ssh_private_key_file=~/ssh/ec2-key.pem

[workers]
# Inputted my Worker Node Public IPs
k8s-worker-frontend ansible_host=3.8.216.19 ansible_user=ubuntu ansible_ssh_private_key_file=~/ssh/ec2-key.pem
k8s-worker-backend ansible_host=35.178.102.235 ansible_user=ubuntu ansible_ssh_private_key_file=~/ssh/ec2-key.pem

[all:vars]
ansible_python_interpreter=/usr/bin/python3
```

Create a playbook to setup kubernetes cluster on the EC2 instances.

```
---
- name: Kubernetes Common Configuration
  hosts: all
  become: yes
  tasks:
```

```

- name: Disable swap
  shell: |
    swapoff -a
    sed -i '/ swap / s/^/#/' /etc/fstab

- name: Load kernel modules
  copy:
    content: |
      overlay
      br_netfilter
  dest: /etc/modules-load.d/k8s.conf

- name: Modprobe modules
  shell: |
    modprobe overlay
    modprobe br_netfilter

- name: Apply sysctl params
  copy:
    content: |
      net.bridge.bridge-nf-call-iptables  = 1
      net.bridge.bridge-nf-call-ip6tables = 1
      net.ipv4.ip_forward                = 1
  dest: /etc/sysctl.d/k8s.conf

- name: Apply sysctl
  command: sysctl --system

- name: Install dependencies
  apt:
    name:
      - ca-certificates
      - curl
      - gnupg
      - lsb-release
      - apt-transport-https
    state: present
    update_cache: yes

- name: Add Docker GPG key
  shell: |
    mkdir -p /etc/apt/keyrings
    curl -fsSL https://download.docker.com/linux/ubuntu/gpg | gpg --dearmor -o
    /etc/apt/keyrings/docker.gpg --yes
    chmod a+r /etc/apt/keyrings/docker.gpg

- name: Add Docker repository
  shell: |
    echo "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg]
https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable" | tee
    /etc/apt/sources.list.d/docker.list > /dev/null

```

```

- name: Install containerd
  apt:
    name: containerd.io
    state: present
    update_cache: yes

- name: Configure containerd
  shell: |
    mkdir -p /etc/containerd
    containerd config default | tee /etc/containerd/config.toml
    sed -i 's/SystemdCgroup = false/SystemdCgroup = true/' /etc/containerd/config.toml
  notify: restart containerd

- name: Add Kubernetes GPG key
  shell: |
    curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.30/deb/Release.key | gpg --dearmor -o
/etc/apt/keyrings/kubernetes-apt-keyring.gpg --yes

- name: Add Kubernetes repository
  shell: |
    echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.30/deb/ ' | tee /etc/apt/sources.list.d/kubernetes.list

- name: Install Kubernetes components
  apt:
    name:
      - kubelet
      - kubeadm
      - kubectl
    state: present
    update_cache: yes

- name: Hold Kubernetes packages
  shell: apt-mark hold kubelet kubeadm kubectl

handlers:
- name: restart containerd
  service:
    name: containerd
    state: restarted
    enabled: yes

- name: Master Node Setup
  hosts: masters
  become: yes
  tasks:
    - name: Initialize Kubernetes Control Plane
      command: kubeadm init --pod-network-cidr=192.168.0.0/16
      register: kubeadm_output
      ignore_errors: yes # Ignore if already initialized

    - name: Setup kubeconfig for ubuntu user

```

```

shell: |
  mkdir -p /home/ubuntu/.kube
  cp /etc/kubernetes/admin.conf /home/ubuntu/.kube/config
  chown $(id -u ubuntu):$(id -g ubuntu) /home/ubuntu/.kube/config

- name: Install Calico Network Plugin
  command: kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml --
kubeconfig=/etc/kubernetes/admin.conf

- name: Generate join command
  command: kubeadm token create --print-join-command
  register: join_command_raw

- name: Set join command fact
  set_fact:
    join_command: "{{ join_command_raw.stdout }}"

- name: Save join command to local file (for debugging/reference)
  local_action: copy content="{{ join_command_raw.stdout }}" dest=./join_command.sh
  become: no

- name: Add dummy host to store join command for workers
  add_host:
    name: "K8S_TOKEN HOLDER"
    join_command: "{{ join_command_raw.stdout }}"

- name: Worker Node Setup
  hosts: workers
  become: yes
  tasks:
    - name: Join Cluster
      shell: "{{ hostvars['K8S_TOKEN HOLDER']['join_command'] }}"
      args:
        executable: /bin/bash
      when: hostvars['K8S_TOKEN HOLDER'] is defined

```

2. Run the playbook `ANSIBLE_HOST_KEY_CHECKING=False ansible-playbook -i inventory.ini playbook.yml`

```

TASK [Setup kubeconfig for ubuntu user] ****
changed: [k8s-master]

TASK [Install Calico Network Plugin] ****
changed: [k8s-master]

TASK [Generate join command] ****
changed: [k8s-master]

TASK [Set join command fact] ****
ok: [k8s-master]

TASK [Save join command to local file (for debugging/reference)] ****
changed: [k8s-master -> localhost]

TASK [Add dummy host to store join command for workers] ****
changed: [k8s-master]

PLAY [Worker Node Setup] ****

TASK [Gathering Facts] ****
ok: [k8s-worker-backend]
ok: [k8s-worker-frontend]

TASK [Join Cluster] ****
changed: [k8s-worker-backend]
changed: [k8s-worker-frontend]

PLAY RECAP ****
k8s-master : ok=24    changed=16   unreachable=0    failed=0    skipped=0    rescued=0    ignored=1
k8s-worker-backend : ok=18    changed=11   unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
k8s-worker-frontend : ok=18    changed=11   unreachable=0    failed=0    skipped=0    rescued=0    ignored=0

kris@Christian:/mnt/c/Users/KRIS/Desktop/t4dev-capstone/infrastructure/ansible$ 

```

3. Check the status of the cluster created in ec2 kubectl get nodes

```

ubuntu@master-node:~$ kubectl get nodes -o wide
NAME           STATUS  ROLES      AGE     VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE          KERNEL-VERSION   CONTAINER-RUNTIME
k8s-worker-backend   Ready   <none>    2m56s   v1.30.14  10.0.1.94    <none>        Ubuntu 24.04.3 LTS  6.14.0-1018-aws  containerd://2.2.1
k8s-worker-frontend  Ready   <none>    2m56s   v1.30.14  10.0.1.129   <none>        Ubuntu 24.04.3 LTS  6.14.0-1018-aws  containerd://2.2.1
master-node       Ready   control-plane  30m    v1.30.14  10.0.1.183   <none>        Ubuntu 24.04.3 LTS  6.14.0-1018-aws  containerd://2.2.1
ubuntu@master-node:~$ 

```

NOTE: I terminate the EC2 instance and use AWS EKS Cluster

4. AWS EKS Cluster

```

KRIS@Christian MINGW64 ~/Downloads
$ kubectl get nodes -o wide
NAME           STATUS  ROLES      AGE     VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE          KERNEL-VERSION   CONTAINER-RUNTIME
ip-10-0-1-111.eu-west-2.compute.internal  Ready   <none>    11h    v1.35.0-eks-78ce843  10.0.1.111  18.171.241.54  Amazon Linux 2023.10.20260202  6.12.66-88.122.amzn2023.x86_64  containerd://2.1.5
ip-10-0-2-12.eu-west-2.compute.internal   Ready   <none>    11h    v1.35.0-eks-78ce843  10.0.2.12   3.8.187.102   Amazon Linux 2023.10.20260202  6.12.66-88.122.amzn2023.x86_64  containerd://2.1.5
KRIS@Christian MINGW64 ~/Downloads
$ 

```

The screenshot shows the AWS EKS console interface. On the left, there's a sidebar with navigation links for Amazon Elastic Kubernetes Service, Settings, Amazon EKS Anywhere, and Related services like Amazon ECR and AWS Batch. The main content area is titled 'my-eks-cluster10'. It has tabs for Overview, Resources, Compute (which is selected), Networking, Add-ons, Capabilities, Access, Observability, Update history & backups, and Tags. Under the Compute tab, there are sections for Cluster info, Nodes (2), and Node groups (1). The Nodes section lists two nodes: 'ip-10-0-1-111.eu-west-2.compute.internal' and 'ip-10-0-2-12.eu-west-2.compute.internal', both of which are t3.medium instances managed by 'my-nodegroup10' and are ready. The Node groups section shows one group named 'my-nodegroup10' with a desired size of 2.

On my GitHub, I created a repository named ShopMicro-Production and added my kube config data and AWS access key and secret key as secrets in the repository settings.

On ks8 master node i run `cat ~/.kube/config` to get the kube config data and paste it in the repository settings secrets.

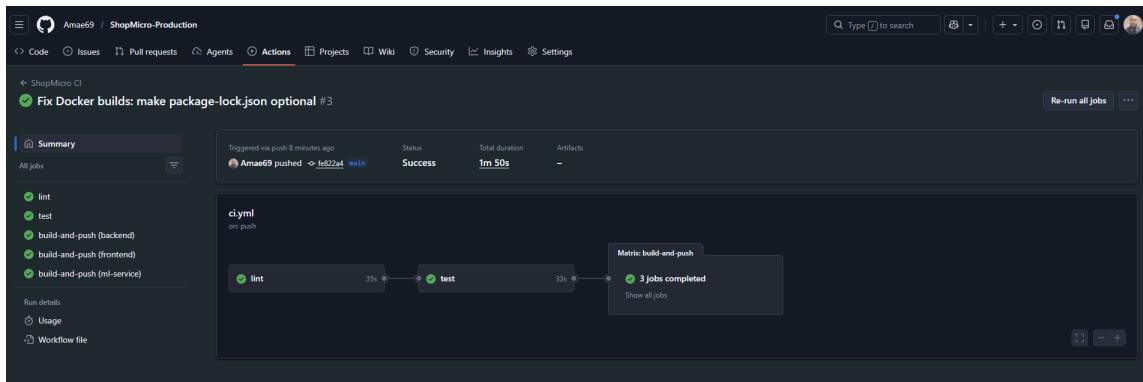
The screenshot shows the GitHub repository settings for 'ShopMicro-Production'. The left sidebar includes options like Code, Issues, Pull requests, Agents, Actions, Projects, Wiki, Security, Insights, and Settings. Under the Settings tab, there's a 'Repository secret added.' message. The main content is the 'Actions secrets and variables' section. It has tabs for Secrets (selected) and Variables. Below that is the 'Environment secrets' section, which says 'This environment has no secrets.' and has a 'Manage environment secrets' button. At the bottom is the 'Repository secrets' table, which lists three secrets: 'AWS_ACCESS_KEY_ID', 'AWS_SECRET_ACCESS_KEY', and 'KUBECONFIG'. Each secret has a 'Last updated' timestamp (4 minutes ago or 5 minutes ago), an edit icon, and a delete icon. A red box highlights this table.

CI/CD Pipeline Implementation

I implemented a complete GitHub Actions workflow suite:

- **ci.yml:**

Triggers on Push/PR. Runs linting (Node/Python), Unit Tests, and builds/pushes Docker images to GHCR.



- **cd.yml:**

Triggers after CI success on main. Deploys the application to my Kubernetes cluster.

The screenshot shows the GitHub Actions interface for a repository named 'ShopMicro-Production'. The pipeline is titled 'ShopMicro CD' and has reached step #45. The 'Actions' tab is selected. On the left, there's a sidebar with links for Summary, All jobs, Run details, Usage, and Workflow file. The main area displays a list of steps for the 'deploy-dev' run, which succeeded 2 minutes ago in 1m 52s. The steps listed are:

- > ✓ Set up job
- > ✓ Checkout Code
- > ✓ Configure AWS Credentials
- > ✓ Update Kubeconfig
- > ✓ Verify Cluster Connection
- > ✓ Update Image Tags
- > ✓ Deploy Namespace
- > ✓ Deploy Secrets
- > ✓ Deploy Postgres
- > ✓ Deploy Backend
- > ✓ Deploy ML Service
- > ✓ Deploy Frontend
- > ✓ Deploy Observability & Other Services
- > ✓ Post Configure AWS Credentials
- > ✓ Post Checkout Code
- > ✓ Complete job

```
kubectl get pod -n shopmicro to check running pod on cluster
```

```
KRIS@Christian MINGW64 ~/downloads
$ kubectl get pods -n shopmicro
NAME                               READY   STATUS    RESTARTS   AGE
backend-5c6fffbc4-66zb4           1/1     Running   0          8m57s
backend-5c6fffbc4-c2dv2           1/1     Running   0          8m42s
frontend-5485bd4b4b-qmg6x         1/1     Running   0          8m10s
grafana-cbfdc9479-6nwts          1/1     Running   0          80m
loki-57fcfb64b4-g4fnn            1/1     Running   0          7m55s
ml-service-deployment-85f59859d9-4tklw 1/1     Running   0          8m27s
postgres-0                         1/1     Running   0          155m
prometheus-c4774cccf-b7zjn        1/1     Running   0          80m
redis-9ccf5476c-4jtm7             1/1     Running   0          80m
tempo-d778cfb78-pflc8             1/1     Running   0          7m48s
```

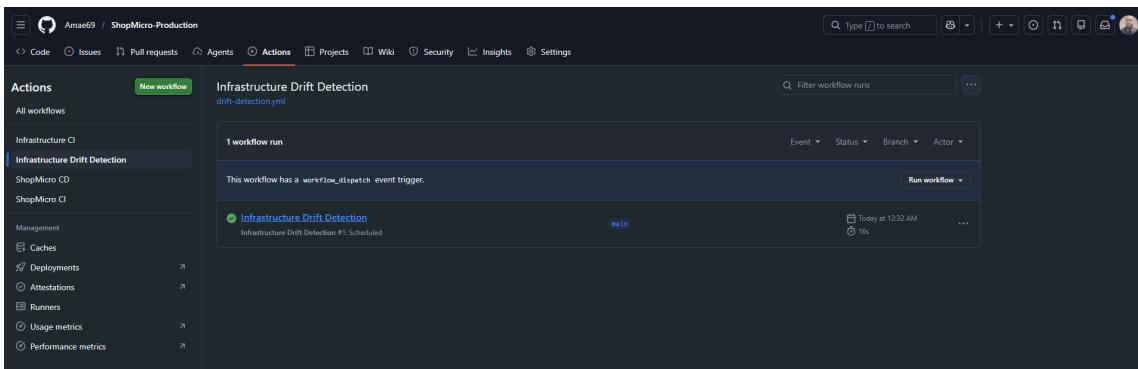
```
KRIS@Christian MINGW64 ~/downloads
$
```

- **iac-ci.yml:**

Runs terraform fmt, terraform validate, tflint, and OPA policy checks on infrastructure changes.

- **drift-detection.yml:**

Runs daily at 8am to check for infrastructure drift. Status: ✓ Implemented. Requires Secrets setup in GitHub.



Configure ingress

Install NGINX Ingress Controller:

The Ingress resource (**k8s/ingress.yaml**) requires a controller to work. Install it on EKS

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/main/deploy/static/provider/aws/deploy.yaml
```

confirm: `kubectl get pods -n ingress-nginx`

```
KRIS@Christian MINGW64 ~/downloads
$ kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/main/deploy/static/provider/aws/deploy.yaml
namespace/ingress-nginx created
serviceaccount/ingress-nginx created
serviceaccount/ingress-nginx-admission created
role.rbac.authorization.k8s.io/ingress-nginx created
role.rbac.authorization.k8s.io/ingress-nginx-admission created
clusterrole.rbac.authorization.k8s.io/ingress-nginx created
clusterrole.rbac.authorization.k8s.io/ingress-nginx-admission created
rolebinding.rbac.authorization.k8s.io/ingress-nginx created
rolebinding.rbac.authorization.k8s.io/ingress-nginx-admission created
clusterrolebinding.rbac.authorization.k8s.io/ingress-nginx created
clusterrolebinding.rbac.authorization.k8s.io/ingress-nginx-admission created
configmap/ingress-nginx-controller created
service/ingress-nginx-controller created
service/ingress-nginx-controller-admission created
deployment.apps/ingress-nginx-controller created
job.batch/ingress-nginx-admission-create created
job.batch/ingress-nginx-admission-patch created
ingressclass.networking.k8s.io/nginx created
validatingwebhookconfiguration.admissionregistration.k8s.io/ingress-nginx-admission created

KRIS@Christian MINGW64 ~/downloads
$ kubectl get pods -n ingress-nginx
NAME                      READY   STATUS    RESTARTS   AGE
ingress-nginx-controller-7fdf8d9764-gsq69   1/1     Running   0          25s

KRIS@Christian MINGW64 ~/downloads
$
```

Apply Ingress Resource: Apply the ingress.yaml file to expose the application externally.

confirm: kubectl get ingress -n shopmicro

```
KRIS@Christian MINGW64 ~/downloads
$ kubectl get ingress -n shopmicro
NAME           CLASS      HOSTS          ADDRESS                                     PORTS   AGE
shopmicro-ingress   nginx    shopmicro.local  a8d8611eb84e74f7f883a0297dd32158-ebd7925954c29321.elb.eu-west-2.amazonaws.com   80      124m

KRIS@Christian MINGW64 ~/downloads
$ |
```

Test the frontend is working through ingress:

```
curl -H "Host: shopmicro.local" http://a8d8611eb84e74f7f883a0297dd32158-
ebd7925954c29321.elb.eu-west-2.amazonaws.com
```

```
KRIS@Christian MINGW64 ~/downloads
$ curl -H "Host: shopmicro.local" http://a8d8611eb84e74f7f883a0297dd32158-ebd7925954c29321.elb.eu-west-2.amazonaws.com
<!doctype html>
<html lang="en">
  <head>
    <meta charset="UTF-8" />
    <meta name="viewport" content="width=device-width, initial-scale=1.0" />
    <title>ShopMicro</title>
    <script type="module" crossorigin src="/assets/index-DF717NJR.js"></script>
  </head>
  <body>
    <div id="root"></div>
  </body>
</html>

KRIS@Christian MINGW64 ~/downloads
$ |
```

Test the backend is working through ingress

```
curl -H "Host: shopmicro.local" http://a8d8611eb84e74f7f883a0297dd32158-
ebd7925954c29321.elb.eu-west-2.amazonaws.com/api/health
```

```
KRIS@Christian MINGW64 ~/downloads
$ curl -H "Host: shopmicro.local" http://a8d8611eb84e74f7f883a0297dd32158-ebd7925954c29321.elb.eu-west-2.amazonaws.com/api/health
{"status": "ok", "service": "backend"}
KRIS@Christian MINGW64 ~/downloads
$
```

Horizontal Pod Auto Scaling:

The HPA is now deployed. Check its status (requires `metrics-server`):

NOTE: if the cpu, or memory threshold is reached, pod will automatically scaled till it get to its set maxpods

```
kubectl get hpa -n shopmicro
```

```
KRIS@Christian MINGW64 ~/downloads
$ kubectl get hpa -n shopmicro
NAME          REFERENCE           TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
backend-hpa   Deployment/backend   cpu: 2%/70%, memory: 50%/80%   2          5          2          129m
ml-service-hpa Deployment/ml-service-deployment   cpu: 2%/60%           1          5          1          129m
KRIS@Christian MINGW64 ~/downloads
$
```

Monitoring (Prometheus & Grafana)

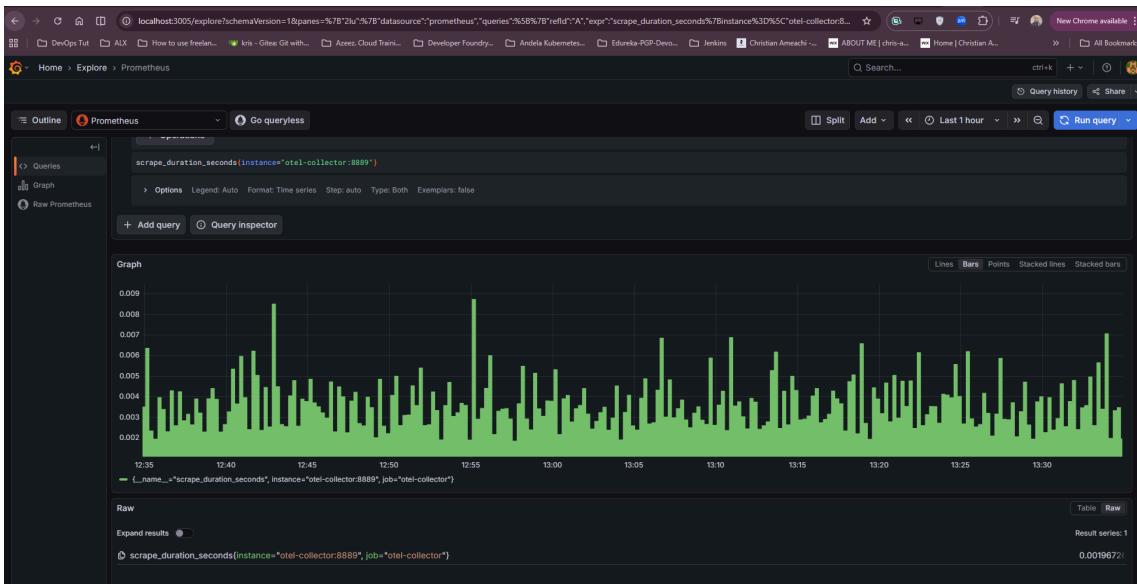
Using portforwarding to test prometheus and grafan access on the web, given that their services are clusterIP and can only be access internally

Portforward prometheus: `kubectl port-forward svc/prometheus 9090:9090 -n shopmicro`

On browser: <http://localhost:9090>

Portforward grafana: `kubectl port-forward svc/grafana 3005:3000 -n shopmicro`

On browser: <http://localhost:3005>



Rollback Proof Demonstration

To demonstrate a successful rollback, I performed the following steps to simulate a failed deployment and recover:

1. Deploy a "Broken" Version:

Manually update the backend to use a non-existent image tag. This will trigger a rolling update that fails (`ImagePullBackOff`).

```
kubectl set image deployment/backend backend=ghcr.io/amae69/shopmicro-production-backend:broken-v1 -n shopmicro
```

2. Observe the Failure:

Check the rollout status. You will see it hang or show errors.

```
kubectl rollout status deployment/backend -n shopmicro  
# Pods will likely show 'ImagePullBackOff'  
kubectl get pods -n shopmicro
```

Screenshot showing "Evidence of Failure".

```
KRIS@Christian MINGW64 ~/downloads
$ kubectl set image deployment/backend backend=ghcr.io/amae69/shopmicro-production-backend:broken-v1 -n shopmicro
deployment.apps/backend image updated

KRIS@Christian MINGW64 ~/downloads
$ kubectl rollout status deployment/backend -n shopmicro
Waiting for deployment "backend" rollout to finish: 1 out of 2 new replicas have been updated...

KRIS@Christian MINGW64 ~/downloads
$ kubectl get pods -n shopmicro
NAME          READY   STATUS        RESTARTS   AGE
backend-76d8b674c8-zpj2k  0/1    ImagePullBackOff  0          7m33s
backend-cb58485fd-d4slt  1/1    Running       0          168m
backend-cb58485fd-dl58c  1/1    Running       0          168m
frontend-77fd755987-499w7 1/1    Running       0          168m
grafana-cbfdc9479-6nwts  1/1    Running       0          5h27m
loki-57fcfb64b4-g4fnn   1/1    Running       0          4h15m
ml-service-deployment-648bfff979-hhl7t 1/1    Running       0          168m
postgres-0              1/1    Running       0          6h42m
prometheus-c4774cccf-b7zjn 1/1    Running       0          5h27m
redis-9ccf5476c-4jtm7   1/1    Running       0          5h27m
tempo-d778cfb78-pflc8   1/1    Running       0          4h15m

KRIS@Christian MINGW64 ~/downloads
$ |
```

3. Perform the Rollback:

Use the `undo` command to revert to the last stable state.

```
kubectl rollout undo deployment/backend -n shopmicro
```

4. Verify Restoration:

Confirm the deployment is back to a healthy state.

```
kubectl rollout status deployment/backend -n shopmicro
kubectl get pods -n shopmicro
```

Screenshot showing "Rollback Proof".

```
KRIS@Christian MINGW64 ~/downloads
$ kubectl rollout undo deployment/backend -n shopmicro
deployment.apps/backend rolled back

KRIS@Christian MINGW64 ~/downloads
$ kubectl rollout status deployment/backend -n shopmicro
deployment "backend" successfully rolled out

KRIS@Christian MINGW64 ~/downloads
$ kubectl get pods -n shopmicro
NAME          READY   STATUS        RESTARTS   AGE
backend-cb58485fd-d4slt  1/1    Running       0          174m
backend-cb58485fd-dl58c  1/1    Running       0          174m
frontend-77fd755987-499w7 1/1    Running       0          173m
grafana-cbfdc9479-6nwts  1/1    Running       0          5h33m
loki-57fcfb64b4-g4fnn   1/1    Running       0          4h20m
ml-service-deployment-648bfff979-hhl7t 1/1    Running       0          174m
postgres-0              1/1    Running       0          6h48m
prometheus-c4774cccf-b7zjn 1/1    Running       0          5h33m
redis-9ccf5476c-4jtm7   1/1    Running       0          5h33m
tempo-d778cfb78-pflc8   1/1    Running       0          4h20m

KRIS@Christian MINGW64 ~/downloads
$
```

END ...